

Displaying large data

Hadley Wickham

Assistant Professor / Dobelman Family Junior Chair
Department of Statistics / Rice University

February 2010



Friday, 12 February 2010

1. Introduction to the diamonds data
2. Histograms and bar charts
3. More boxplots
4. Scatterplots for large data
5. Some theory

Diamonds data

~**54,000** round diamonds from
<http://www.diamondse.info/>

Carat, colour, clarity, cut


Total depth, table, depth,
width, height

Price



← x →

← table width →



$$\text{depth} = z / \text{diameter}$$
$$\text{table} = \text{table width} / x * 100$$

Histogram & bar charts

Histograms and barcharts

Used to display the **distribution** of a
variable

Categorical variable → bar chart

Continuous variable → histogram

Always
experiment with
the bin width!

Examples

```
# With only one variable, qplot guesses that  
# you want a bar chart or histogram  
qplot(cut, data = diamonds)
```

```
qplot(carat, data = diamonds)  
qplot(carat, data = diamonds, binwidth = 1)  
qplot(carat, data = diamonds, binwidth = 0.1)  
qplot(carat, data = diamonds, binwidth = 0.01)  
resolution(diamonds$carat)
```

```
last_plot() + xlim(0, 3)
```


Examples

```
# With only one variable, qplot guesses that  
# you want a bar chart or histogram  
qplot(cut, data = diamonds)
```

```
qplot(carat, data = diamonds)  
qplot(carat, data = diamonds, binwidth = 1)  
qplot(carat, data = diamonds, binwidth = 0.1)  
qplot(carat, data = diamonds, binwidth = 0.01)  
resolution(diamonds$carat)
```

Common ggplot2
technique: adding
together plot
components

```
last_plot() + xlim(0, 3)
```

```
qplot(table, data = diamonds, binwidth = 1)


# To zoom in on a plot region use xlim() and ylim()
qplot(table, data = diamonds, binwidth = 1) +
  xlim(50, 70)
qplot(table, data = diamonds, binwidth = 0.1) +
  xlim(50, 70)
qplot(table, data = diamonds, binwidth = 0.1) +
  xlim(50, 70) + ylim(0, 50)

# Note that this type of zooming discards data
# outside of the plot regions
# See coord_cartesian() for an alternative
```


Additional variables


As with scatterplots can use **aesthetics** or **faceting**. Using aesthetics creates pretty, but ineffective, plots.


The following examples show the difference, when investigation the relationship between cut and depth.



```
qplot(depth, data = diamonds, binwidth = 0.2)
```







```
qplot(depth, data = diamonds, binwidth = 0.2) +
  xlim(55, 70) + facet_wrap(~cut)
```

Your turn

Explore the distribution of price.

How does it vary with colour, or cut, and clarity?

Weighting

```
qplot(cut, data = diamonds, weight = carat)  
qplot(cut, data = diamonds, weight = price)
```


```
# Also useful for pretabulated data  
cuts <- as.data.frame(table(  
  cut = diamonds$cut))  
qplot(cut, weight = Freq, data = cuts)
```

Scatterplots


Revision:


Interpreting a scatterplot


- Big patterns
- Small patterns
- Deviations from the pattern
- Strange patterns





Strong linear relationship.
A number of outliers.















But what's the
problem with
all these plots?

```
qplot(carat, price, data = diamonds)
```



Ideas

If x discrete, use boxplots.

Use semi-transparent points.

Divide into bins and count number of points in each bin (2d histogram).


Display statistical summary.

Box and whisker plots


Boxplots

Less information than a histogram, but take up much less space.


Already seen them used with discrete x values. Can also use with continuous x values, by specifying how we want the data **grouped**.




```
qplot(table, price, data = diamonds)
```




```
qplot(table, price, data = diamonds, geom = "boxplot")
```


```
qplot(table, price, data = diamonds, geom = "boxplot",  
      group = round(table))
```



Alpha blending




```
qplot(carat, price, data = diamonds, alpha = I(1/10))
```




Must use I() for
constants

```
qplot(carat, price, data = diamonds, alpha = I(1/10))
```



```
qplot(carat, price, data = diamonds, alpha = I(1/50))
```



```
qplot(carat, price, data = diamonds, alpha = I(1/250))
```


```
qplot(carat, price, data = diamonds,  
      colour = I(alpha("black", 1/255)))
```

```
qplot(carat, price, data = diamonds, geom = "bin2d")  
qplot(carat, price, data = diamonds, geom = "bin2d",  
      bins = 100)
```


```
qplot(carat, price, data = diamonds, geom = "hex")
```

```
qplot(carat, price, data = diamonds) + geom_smooth()
```



Statistical summary



```
qplot(carat, price, data = diamonds) + geom_smooth()
```



```
qplot(log10(carat), log10(price), data = diamonds) + geom_smooth()
```



```
qplot(log10(carat), log10(price), data = diamonds) +  
  geom_smooth(method = "lm")
```

2d bins

```
# Very basic cleaning
```

```
diamonds$x[diamonds$x == 0] <- NA
```

```
diamonds$y[diamonds$y == 0] <- NA
```

```
diamonds$y[diamonds$y > 12] <- NA
```

```
qplot(x, y, data = diamonds)
```

```
qplot(x, y, data = diamonds, geom = "bin2d")
```

```
qplot(x, y, data = diamonds, geom = "hex")
```

```
qplot(x, y, data = diamonds, geom = "bin2d", bins = 100)
```

```
qplot(x, y, data = diamonds, geom = "hex", bins = 100)
```

```
# Zoom in
```

```
qplot(x, y, data = diamonds, geom = "bin2d", bins = 100) +  
  xlim(4,7) + ylim(4,7)
```

```
qplot(x, y, data = diamonds, geom = "bin2d", bins = 100) +  
  xlim(4,5) + ylim(4,5)
```

```
qplot(x, x / y, data = diamonds,  
      geom = "bin2d")  
qplot(x, log(x / y), data = diamonds,  
      geom = "bin2d")  
  
clean <- subset(diamonds, abs(log(x / y)) < 0.1)  
  
qplot(x, log(x / y), data = clean, geom = "bin2d")  
qplot(x, log(x / y), data = clean, geom = "bin2d",  
      bins = 80)
```

```
qplot(x, x / y, data = diamonds,  
      geom = "bin2d")
```

```
qplot(x, log(x / y), data = diamonds,  
      geom = "bin2d")
```

```
clean <- subset(diamonds, abs(log(x / y)) < 0.1)
```

```
qplot(x, log(x / y), data = clean, geom = "bin2d")  
qplot(x, log(x / y), data = clean, geom = "bin2d",  
      bins = 80)
```

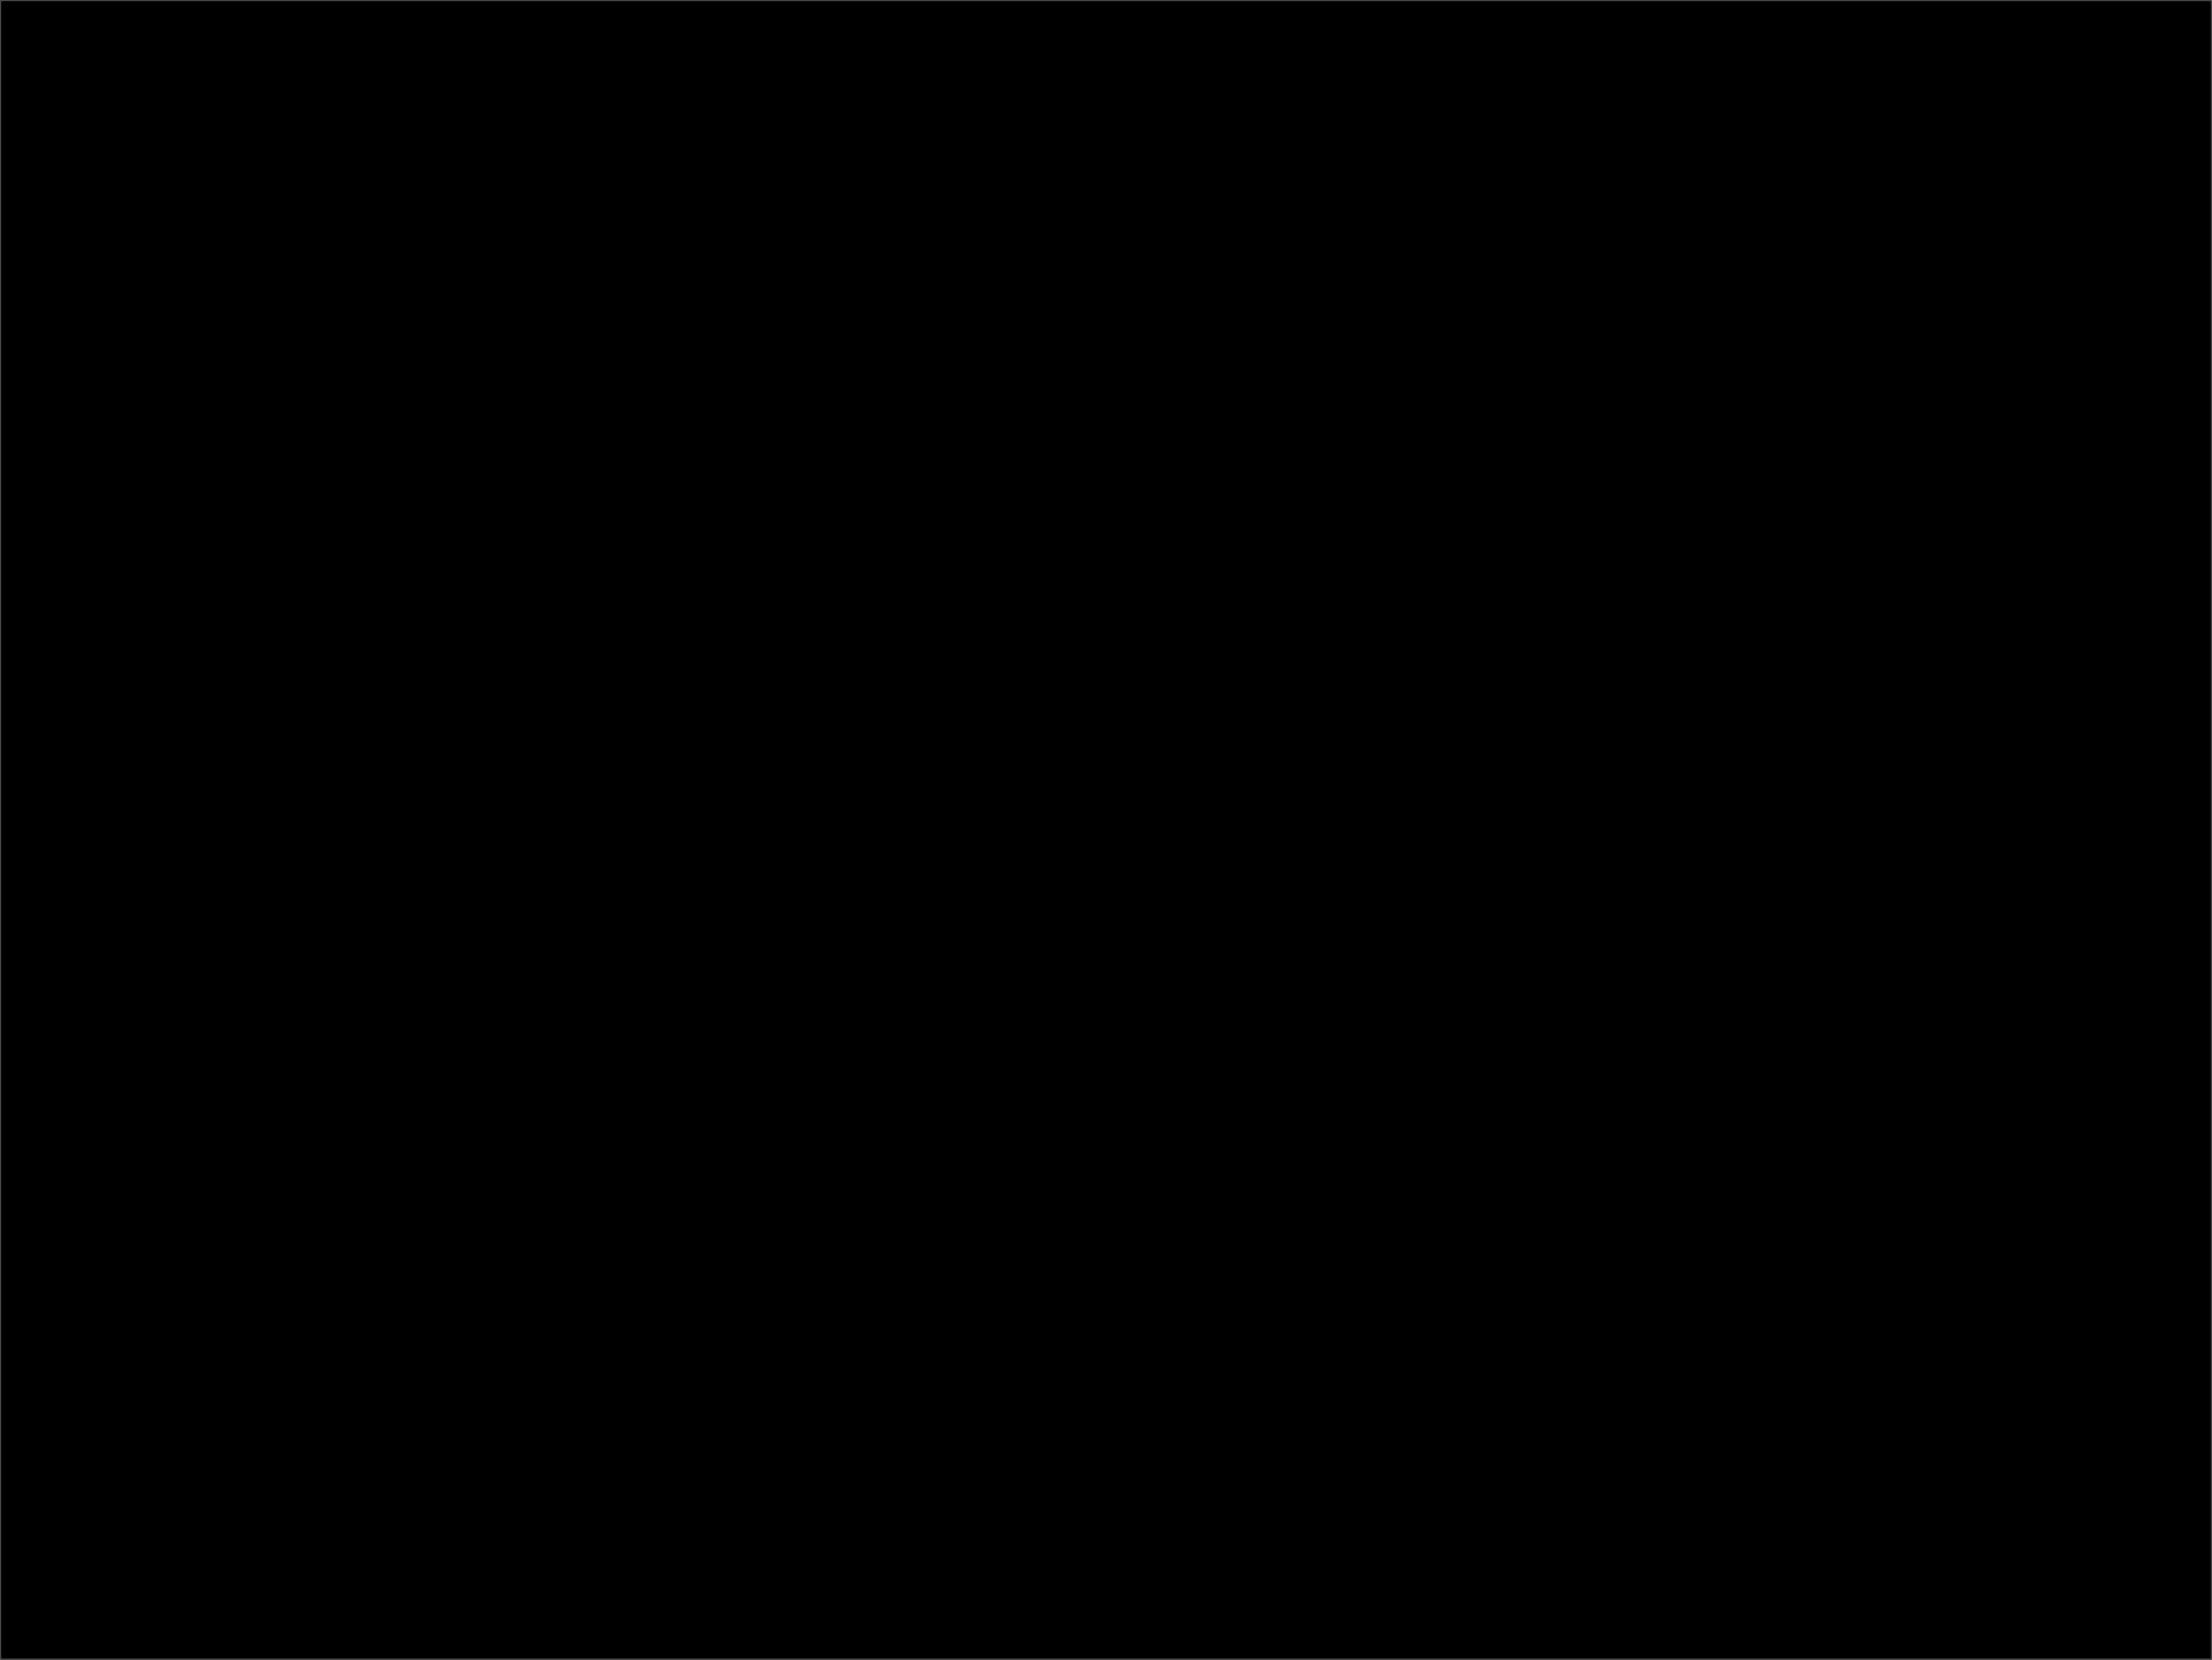
What would be a good name for $\log(x / y)$? What other variable might you create to go with it?

Your turn

Continue to explore the relationship between x, y, z and carat. Create new variables as necessary.

You might also want to do more cleaning.

Some good ideas here: <http://www.diamondhelpers.com/fivesteps/4-certified-diamonds.shtml>



This work is licensed under the Creative Commons Attribution-Noncommercial 3.0 United States License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nc/3.0/us/> or send a letter to Creative Commons, 171 Second Street, Suite 300, San Francisco, California, 94105, USA.